Introduction to Logic Programming in C++

Roshan Naik (roshan@mpprogramming.com)

[DRAFT] Last updated: Oct 22nd, 2008

Version History

Feb 11th 2008: Initial version.
Mar 15th 2008:
- Renamed `GenerativeRelation` to `OneSolutionRelation` throughout.
- Moved out section on "Implementing Relations imperatively" from section 2 into its own top level section 3.
- Rewrite of section 3 in progress. Now described in terms of class Coroutine and co_* macros .

TABLE OF CONTENTS

Abstract:

This paper is an introductory tutorial for Logic paradigm (LP) in C++. No prior experience is required with languages that natively support LP. It also demonstrates how LP blends with the other paradigms supported by C++ and also the STL. The ability to choose an appropriate paradigm or an appropriate mix of paradigms for solving a given problem is essential and at the heart of multiparadigm programming. We begin with a brief introduction to the logic paradigm, followed by a discussion of logic style programming in C++ and finally conclude with examples. The primitives used here for logic programming are provided by Castor, an open source C++ library available from www.mpprogramming.com. No language extensions to C++ are required to compile the code provided here.

1 The Logic paradigm

Logic programming is a Turing-complete programming paradigm. The model of computation used in logic is strikingly different from the more mainstream imperative and functional paradigms. Pure logic programs are entirely declarative in nature. When programming in languages based on the imperative paradigm (like C, C++, Java etc), programmers actively instruct the computer *how* to solve a specific problem and the computer itself has no knowledge about the problem. Thus, algorithms play a central role. In logic programming languages such as Prolog or Gödel, however, it is exactly the opposite. In LP, the programmer deals with providing problem specific information to the computer instead of providing the steps required to solve a specific problem. The computer applies a general purpose problem solving algorithm to the domain specific information to produce the desired results. The programmer is not involved with specifying the exact steps (i.e. the algorithm) used in solving the problem.

Information provided to the computer in logic programs can be classified into *facts* and *rules*. This knowledge-base of facts and rules describes the problem domain. Specific problems that we wish to solve in this domain are posed as questions or *queries*. The computer examines the query in the context of the rules and facts and determines the solution. For example, if the game of Chess (or some other board game) represents our problem domain, the facts may comprise of things such as:
- The different kinds of pieces (e.g. white pawns, black pawns, white king, etc.)
- The number of pieces of each kind (e.g. 8 black pawns, 1 white king, etc.)
- A description of the number of squares and their layout on the board (e.g. 8x8 board, 32 white squares, 32 black squares etc.)

And the rules may comprise of:
- The rule governing the movement of each kind of piece on the board(e.g. bishop moves diagonally)
- The rule to determine if a piece is under attack
- The rule to determine when a game is over and the result of the game.

Typical questions that arise in Chess (and many other board games):
- Given a specific layout of pieces on the board, what are all the possible moves for a given piece?
- Given a specific board layout, which pieces can be moved next?
- Which pieces are under attack in a given board layout?

Each question above represents a different but concrete problem that belongs to the problem domain of Chess. Shifting focus from describing how to solve a particular type of problem to describing the general rules of the broader problem domain, allows us to seek answers to a wider variety of problems within the domain. In the remainder of this section we will further illustrate facts, rules and queries using a simple example concerning family relationships. The primary focus here is to get a feel for the basic mechanics of LP.

1.1. Facts

Facts are essentially the simplest form of true statements pertaining to a problem domain. They may also be referred to as data. Let us consider a four person family (Son: Sam, Daughter: Denise, Father: Frank, Mother: Mary, Grandparent: Garry). Here is one way of describing the facts pertaining to this family more accurately:

Children facts:
1. Sam is a child of Mary
2. Denise is a child of Mary
3. Sam is a child of Frank
4. Denise is a child of Frank
5. Frank is a child of Garry

Gender facts
6. Frank is male
7. Sam is male
8. Mary is female
9. Denise is female

The above facts can be used to answer simple questions such as "Is Sam male?". This is a basic true/false question and can be answered by inspecting the gender facts. Since we have an exact match with fact 7, the answer to the question is "yes" or "true". However, a similar question "Is Frank female?" yields "false" or "no". This is because we do not have any gender fact stating Frank's gender to be female. Thus, whenever matching evidence is found the answer is true; and false otherwise.

A slightly different type of question is "What is the gender of Sam?". This is not a true/false question but still can be answered by looking up the gender fact 7. Another question could be "Who is the child of Frank?". Again this is not a true/false question, however it is a little more interesting as there is not one but two answers to it, *Sam* and *Denise*. This tells us that we cannot simply stop examining the facts as soon

as the first match is found; we need to continue the search till all relevant facts are exhausted. When there are multiple answers to a question, the person asking the question may be interested in one, some or all the answers.

A question that cannot be answered based solely based on the above facts is "Is Mary a parent of Sam?". This is because we have not yet declared what it means to be a *parent*. There are no direct facts stating any parent-of relationships. To solve this we can add one fact of the nature "X is parent of Y" for each fact of the nature "Y is child of X" above. This approach is cumbersome and error-prone especially when the database of facts is large. A better approach is to use *rules* to infer these new facts.

1.2. Rules

Rules are statements used to *infer* new facts (or data) from an existing body of facts. Here are some simple rules pertaining to the family relationships:

Parent rule: X is parent of Y if
1. Y is child of X.

Father rule: X is father of Y if
2. Gender of X is male **and** Y is child of X

Mother rule: X is mother of Y if
3. Gender of X is female **and** X is parent of Y

Here X and Y are essentially parameters and not constants like "Sam" and "Frank". The *parent* rule provides the ability to answer questions like "Is Mary a parent of Sam?" or "Who is a parent of Sam?". Note that the *parent* and *father* rules above are specified only in terms of facts. The *mother* rule on the other hand is specified using a combination of facts (*gender*) and rules (*parent*) although it could be specified in same manner as the *father* rule.

1.2.1. Recursive rules

Rules can also be specified in terms of themselves. Consider describing the ancestor relationship as a rule.
Ancestor rule: X is ancestor of Y if
4. X is parent of Y **or**
5. Z is parent of Y **and** X is ancestor of Z

Now we are in a position to ask "Is Garry an ancestor of Sam?" which should yield *true*. Similarly if we ask "Who is an ancestor of Sam?", it should yield *Frank*, *Mary* and *Garry*.

1.3. Queries and Assertions

Once we have built the knowledge-base consisting of facts and rules we are ready to ask questions. Specific problems that need to be solved are posed as queries. We have seen examples of queries above for the family relationships. Other examples of queries, for instance, when dealing with graphs is to ask "What is the shortest path between nodes A and B?" or "Is graph G a connected graph?".

Queries can be classified into two categories. The first kind simply tests if a certain statement is true: "Is Sam child of Garry?". These are also traditionally referred to as assertions, since the purpose is to essentially check or assert whether a fact is true. The second type of query is one that seeks for one or more solutions. For example if we ask "Who is the child of Frank?", we are not asserting if a fact is true, but instead asking the system to determine Frank's child. We will henceforth refer to these as generative queries as it requires generation of solutions.

1.4. Computation by inferencing

We have not been very specific, so far, about how answers are actually computed (or inferred). Given the facts and rules described above it is fairly intuitive for any individual to mentally infer answers to all the questions we have asked. The fundamental principle behind such inferencing of new facts from existing rules and facts is referred to in formal Logic as *modus ponens:*

> *If A is true and*
> *A implies B,*
> *then B is true.*

A and B, above, refer to arbitrary statements. This principle is at the heart of the computational model used in logic programming. An intuitive but rather rough approximation of the algorithm used to solve queries in logic programming systems such as Prolog is as follows:

1. Build a list of relevant facts and rules
2. Pick a relevant fact and see if it answers the question. Repeat this until all relevant facts are exhausted.
3. Pick a relevant rule that can be applied to derive new facts (using modus ponens). Repeat this until all relevant rules and facts are exhausted.

At any step of the inferencing algorithm, there may be more than one rule and/or fact that can be selected in order to continue execution. Each choice leads to a different inference path and not all paths necessarily lead to solutions. Paths that do not lead to solutions are abandoned and inferencing resumes from the most recent point where other facts and/or rules were available for inferencing. Abandoning the current path and resuming execution from an earlier point in order to make a different choice is known as *backtracking*. Backtracking continues till all paths of inference have been traversed. Thus, computation is reduced to traversing different paths of inference determined from the facts and rules available. This is similar to performing depth first search in a binary tree where leaf nodes represent candidate results and the inner nodes represent intermediate results.

In pure formal logic, the exact order in which facts and rules are selected for application is non-deterministic. This also implies that the order in which answers are obtained is not deterministic. Since some paths may lead to solutions quicker than others, in practice the order of execution is fixed (same as declaration order) to allow control over efficiency. Fixing the order of application of rules also simplifies reasoning about the execution of logic programs which is important for debugging.

1.5.   Summary

In this section we described facts, rules, queries, assertions and inferencing. Facts are simply true statements. Rules can be used to derive new facts. Rules can be specified in terms of facts, other rules or even themselves (i.e. recursive rules). A collection of rules and facts can be used to answer relevant questions. Questions can be broadly classified into ones that simply require a true/false answer or ones for which solutions need to generated. The former types of questions are referred to as assertions and the later referred to as generative queries. Assertions have only one answer (i.e. true/false). Generative queries may have zero or more solutions. "Is Sam male?" is an example of an assertion. "Who is the child of Mary?" is an example of a generative query. Logical inferencing is used to answer questions. It involves examining the facts and application of rules. New facts emerge from application of rules which become candidates for future consideration during inferencing.

2   Logic Programming in C++

Now let us translate the above facts and rules into C++ so that they can be executed. The examples here make use of Castor, an open source library that enables logic programming in C++. Castor enables embedding of logic style code naturally into C++ by allowing rules and facts to be declared as classes, functions or even just expressions. This low level of integration is very useful and allows for a programming platform where the paradigm-boundaries are seamless.

The following C++ functions represent the *child* and *gender* facts and the *father* rule (described previously) using Castor:

```cpp
// c is child of p
relation child(lref<string> c, lref <string> p)
{
  return eq(c,"Sam")    && eq(p,"Mary") //fact 1
    || eq(c,"Denise") && eq(p,"Mary") //fact 2
    || eq(c,"Sam")    && eq(p,"Frank")//fact 3
    || eq(c,"Denise") && eq(p,"Frank")//fact 4
    || eq(c,"Frank")  && eq(p,"Gary") //fact 5
  ;
}

// p's gender is g
relation gender(lref<string> p, lref<string> g)
{
  return eq(p,"Frank")&& eq(g,"male")  //fact 6
    || eq(p,"Sam")    && eq(g,"male")  //fact 7
    || eq(p,"Mary")   && eq(g,"female")//fact 8
    || eq(p,"Denise")&& eq(g,"female")//fact 9
```

```cpp
  ;
}

// f is the father of c
relation father(lref<string> f,lref<string> c)
{
  //... if f is male and c is child of f
  return gender(f,"male")&& child(c,f);//rule2
}
```

Facts and rules are both declared as functions with the return type *relation*. The parameter types are specified in terms of template lref which stands for *logic reference*. Here it provides a facility similar to the pass by reference mechanism in C++. However, unlike references in C++, logic references can be left uninitialized. The value underlying a logic reference can be obtained by dereferencing it with operator *.

Function eq is called the unification relation. Its job is to *attempt* to make its two arguments equal. If any one of its arguments is an uninitialized logic reference, it will assign the other argument to it. If both arguments have well defined values then it simply compares the two arguments. This task is referred to as **unification**. Consider the call eq(c,"Sam") in relation child. If c has been previously initialized with a value, eq will compare the contents of c with "Sam". However if c has not been initialized, "Sam" will be assigned to c. The type relation, logic references and relation eq will be discussed further in sections 2.1, 2.2 and 2.3 respectively.

 Neither eq nor any of the user defined relations like child or gender return the results of their intended computation immediately when invoked. Instead they return function objects that encapsulate the intended computation. The function objects can be stored in an object of type relation their evaluation can be triggered by application of the function call operator on relations. Given the above C++ definitions for the relations we are in a position to make some queries and assertions. The following is a simple assertion to check if Sam is male:

```cpp
relation samIsMale = gender("Sam", "male");
if( samIsMale() )
  cout << "Sam is male";
else
  cout << "Sam is not male";
```

Similarly we can check if Frank is Sam's father:

```cpp
relation samsDadFrank = father("Frank","Sam");
if(samsDadFrank())
    cout << "Frank is Sam's father";
else
    cout << "Frank is not Sam's father";
```

We can also issue generative queries such as "What is Sam's gender?":

```cpp
lref<string> g; // g not initialized
```

```
relation samsGender = gender("Sam", g);
if( samsGender() )
    cout << "Sam's gender is " << *g;
else
    cout << "Sam's gender is not known";
```

Here we pass "Sam" as the first argument and simply leave the second argument undefined (i.e. not initialized to any value). Note how the same function gender is used to assert if Sam is male and also find out Sam's gender. When all the arguments have been defined (i.e. initialized to some value) it performs a check to see if they satisfy the gender relation. Arguments that have been left undefined will act as output parameters and will be initialized with a value that satisfies the gender relation. This bidirectional nature of the lref parameters behaving as input or output is central to the logic programming model. It is important to note however that lref parameters on relations typically do not act as both input and output simultaneously as is often the case when using the pass-by-reference scheme in functions like swap.

What happens if both arguments to gender relation are undefined?

```
lref<string> p, g; // p and g not initialized
relation anyPersonsGender = gender(p, g);
if( anyPersonsGender() )
    cout << *p << "'s gender is " << *g;
```

In this case both p and g will be assigned values by gender. Since the person-gender pair ("Frank", "male") is declared first in the gender relation, p will be assigned "Frank" and g will be assigned "male".

Generative queries, such as samsGender and anyPersonsGender above, may have zero, one or many solutions. So far we have only generated the one solution from them. Iterating over all solutions, quite naturally, involves the use of a while loop instead of the if statement we have used so far. The previous example can be rewritten as follows to print all persons and their genders as follows:

```
while( anyPersonsGender() )
    cout << *p << "'s gender is " << *g << "\n";
```

Similarly for listing of all Frank's children:

```
lref<string> c;
int count=0;
relation franksChildren = father("Frank", c);
while( franksChildren() ) {
    ++count;
    cout << *c << " is Frank's child\n";
}
// c is now back to its uninitialized state
cout << "Frank has " << count << " children";
```

Once all solutions have been exhausted, invoking franksChildren() returns false and causes the while loop

to terminate. Also when all solutions have been exhausted logic reference c will be automatically reset to its original uninitialized state.

Notice how the ability to use fundamental imperative constructs like the if statement and the while loop makes the transition between the logic programming model (using which results are generated) and the imperative model (where the results are consumed) simple and seamless.

Function eq, template type lref and type relation along with overloads for operators && and || provide the foundation for logic programming in Castor. They are described briefly in the following sections. For an in depth coverage of their design and implementation refer to [CastorDesign].

2.1   Type relation

In logic programming it is common to refer to facts and rules as **predicates** or **relations**. The term "relation" originates in set theory where it is used to imply an association between sets. So, gender is a **binary relation** between a set of individuals and the set of genders. Generally a strict distinction between rules and facts is not required when programming with Castor as they can be mixed freely within the same function/expression[1]. Keeping with logic programming lingo, we will henceforth refer to functions that represent facts or rules as **relations**. So, functions with return type relation, are themselves referred to as relations.

The type relation internally represents a function or function object with no arguments and return type bool. Thus child, gender and father return a function object which can be evaluated later in a lazy manner.

2.2   lref: Logic reference

Template type lref is an abbreviation for *logic reference* and provides a facility for passing values in/out of relations in the form of arguments, similar to references in C++. Unlike C++ references, a logic reference does not have to be initialized and provides the member function defined for checking if it has been initialized. The dereference operator * and the arrow operator -> can be applied to access the value referred to by a logic reference.

Now let us understand the initialization semantics of lref semantics. This is helpful reasoning about operational semantics of relational code. When an lref<T> is initialized (i.e. constructed) or assigned a value of type T (or a type

---

[1] This is different from the approach taken in classical logic programming systems like Prolog. The general approach, taken by Castor, of how to blend relations syntactically with the imperative languages was pioneered by Timothy Budd in his multiparadigm language Leda. Although Castor shamelessly steals this idea from Leda, the underlying implementation techniques diverge significantly.

convertible to `T`), it internally stores a copy of the value. When initialized with another `lref<T>` (i.e. copy constructed), both logic references will be *bound together*. References that are bound together will refer to the same underlying value (if any). Thus any change to the underlying value of one logic reference is observed by all logic references that are bound together. A binding between logic references cannot be broken. That is, if logic references *A* and *B* are bound together and *C* and *D* are bound together, then *C*'s binding with *D* cannot be broken in order to form a binding with *A* and *B*. C will continue to be a part of the binding for the duration of its lifetime. Logic references can only be bound by initialization (i.e. copy construction) and not by assignment. A binding can only be formed during construction of the logic reference and will be automatically broken when the logic reference is destroyed. When the last logic reference that is part of the binding is destroyed, it will deallocate the underlying value.

2.3   Relation eq: The unification function

Function `eq` is the unification function and takes two arguments. The arguments may be logic references or regular values. `eq` returns an expression (i.e. function object) which when evaluated attempts to *unify* the two arguments. If unification succeeds it returns true and false otherwise. Unification performed by `eq` is defined as follows.
- If both arguments are initialized, their values are compared for equality and the result of comparison is returned
- If the only one argument is initialized, the uninitialized argument will be is assigned the value of initialized one in order to make them equal.
- If both arguments are uninitialized an exception is thrown.

So unification will *generate* a value for the uninitialized argument to make both arguments equal, or it will *compare* its arguments if both are initialized. In short, unification is a "generate or compare" operation. It is possible to implement other variations to unification, but is generally not required.

The expressions returned by the various calls to `eq` within, for instance, the `gender` relation are stitched together to form a bigger compound expression using ‖ and && operators. It is important to note that the resulting compound expression is returned without being evaluated. These expressions are evaluated in a lazy manner. In other words, these expressions are stored in an object of type `relation` and will be subject to evaluation in the future when needed. The following section examines the evaluation of these expressions in detail.

2.4   Evaluating Queries

Given the above relations, we can formulate the query that tests for "Is Sam male?" and store it in a variable samIsMale as follows:

```
relation samIsMale = gender("Sam", "male");
```

The call `gender("Sam","male")` does not perform any real computation; it simply returns a function object that can be later executed to determine if Sam is male. Invocation of a relation does not execute/evaluate it. This splitting of invocation from execution is a case of *lazy evaluation*. In order to execute this expression we simply apply the function call operator on the variable that holds the expression:

```
if( samIsMale() )
  cout << "Sam is male";
else
  cout << "Sam is not male";
```

The `gender` relation was defined previously as follows:

```
// p's gender is g
relation gender(lref<string> p,lref<string> g)
{
  return eq(p,"Frank")  && eq(g,"male")
      || eq(p,"Sam")    && eq(g,"male")
      || eq(p,"Mary")   && eq(g,"female")
      || eq(p,"Denise") && eq(g,"female");
}
```

Since "Sam" and "male" were passed as arguments to `gender`, the expression returned by `gender` looks as follows with the arguments substituted:

```
   eq("Sam", "Frank")  && eq("male", "male")
|| eq("Sam", "Sam")    && eq("male", "male")
|| eq("Sam", "Mary")   && eq("male", "female")
|| eq("Sam", "Denise") && eq("male", "female")
```

This expression contains four && expressions joined by three ‖ operators. If any one of these && expressions returns `true`, the overall expression has been solved and evaluation halts in order to report success. Let us step through the execution that takes place when **operator()** is applied to samIsMale. The first expression eq("Sam", "Frank") && eq("male", "male") is chosen for evaluation. An attempt to unify "Sam" with "Frank", when evaluating eq("Sam", "Frank"), fails as "Sam" is not equal to "Frank". Short-circuit evaluation principle tells us that the remainder of this && expression does not need to be evaluated. Thus, this path of execution is abandoned immediately and backtracking will resume execution from the next && expression eq("Sam", "Sam") && eq("male", "male"). This time both halves of the && expression unify successfully as their arguments are equal. A solution to the expression has been found and true is returned to the caller which happens to be the if statement.

If the function call operator is applied once again to samIsMale, execution resumes from the point at which it was previously halted. In this case the third && expression is chosen for evaluation, which fails due to failure in unification of "Sam" and "Mary". This leads to evaluation of the fourth and final && expression which also fails for similar reasons

and false is returned to the caller. All `&&` expressions have now been evaluated and applying the function call operator to `samIsMale` henceforth will immediately return false.

Now let us consider how to specify generative queries such as "What is Sam's gender?". Such a query is constructed by calling `gender` with first argument initialized to "Sam" and leaving the second argument uninitialized:

```
lref<string> g; // g not initialized
relation samsGender = gender("Sam", g);
if( samsGender() )
   cout << "Sams gender is " << *g;
```

Notice how it is possible to use the same `gender` relation to both assert if someone is male and to find the gender of a given individual. When a solution is found, `g` will be initialized to "male". The expression returned by `gender("Sam", g)` now looks like this once we substitute the arguments:

```
   eq("Sam","Frank")   && eq(g,"male")
|| eq("Sam","Sam")     && eq(g,"male")
|| eq("Sam","Mary")    && eq(g,"female")
|| eq("Sam","Denise")  && eq(g,"female")
```

When `operator()` is applied to `samsGender` for the first time, the first `&&` expression fails evaluation as "Sam" does not unify with "Frank". Backtracking proceeds to try out the second expression `eq("Sam", "Sam") && eq(g, "male")`. Unification of "Sam" with "Sam" succeeds and then `eq(g, "male")` is evaluated. Since `g` is undefined, `eq` will assign "male" to `g` and thus unification succeeds. Evaluation of the entire expression halts and returns `true` to the calling `if` statement. If `operator()` is again applied to `samsGender`, execution will resume from the point it was previously halted. However, one very interesting thing happens before execution proceeds. It is critical that side effects occurring in the previous path of execution must be undone before attempting a different alternative. If side effects are not reverted before backtracking pursues another alternative, the side effects can affect the results of future evaluations and lead to incorrect results. Thus unification of `g` with "male" has to reverted causing `g` to go back to its original uninitialized state. This undo feature is automatically provided by the unification function `eq`.

The above query has only one solution, but as we observed previously, queries can have multiple solutions. For instance we may want to find all males in the system. Once again we resort to the `gender` relation, but this time we leave the first argument uninitialized and initialize the second argument to "male":

```
  lref<string> person;
  relation males = gender(person, "male");
  while( males() )
      cout << *person << " is male\n";
```

This time around we repeatedly invoke `males()` till it returns `false`. Each invocation of `males` triggers the search for the next solution i.e. a suitable value for the logic reference `person`. Each time around the loop, `person` is assigned a different value representing a solution. When all solutions have been discovered, `males()` returns false, the loop terminates and backtracking will restore `person` to its original uninitialized state. Since `person` does not refer to anything after the while loop has terminated, any attempts to access the underlying value using operator `*` or operator `->` will throw an exception.

2.5   Recursive Rules

Recursion is often essential when declaring rules in the logic paradigm. For the family relationships example, let us consider defining a rule for the "ancestor" relationship. Note that there can be arbitrary levels of parent-child relationships between an ancestor and a descendant. We can recursively define the rule as:

A is an ancestor of D if:
- D is a child of A, **OR**
- D is a child of some P **AND** A is an ancestor of P

This lends itself naturally into the following definition for the ancestor relation

```
// Flawed recursive definition
relation ancestor(lref<string> A
                 , lref<string> D) {
  lref<string> P;
  return child(D, A)
     || child(D, P) && ancestor(A, P);
}
```

However, the above C++ definition contains infinite recursion. The return statement contains a recursive call to the `ancestor`. In order to break the recursion we need to postpone the recursive invocation so that it only takes place if truly needed. Castor provides helper relation `recurse` for defining recursive rules. `recurse` takes the relation that needs to called recursively and the arguments that need to be used for the call. It returns a function object which when evaluated leads to the actual recursive call to `ancestor`. The problematic call `ancestor(A,P)` is simply rewritten as `recurse(ancestor, A, P)`:

```
relation ancestor(lref<string> A
                 , lref<string> D) {
  lref<string> P;
  return child(D,A)
     || child(D,P)
            && recurse(ancestor, A, P);
}
```

In the above example, recursion is performed on relation `ancestor` which is defined as a global function. Usage of `recurse` is the same for static member relations too. For

recursing on non-static member relations, the `this` pointer needs to be provided in addition to the method name and arguments as follows.

```
recurse(this, &Type::method, ..args..)
```

2.6   Dynamic Relations

Definitions of all relations described so far have been fixed at compile time. Relation `gender`, for instance, provides a definite list of name-gender pairs that us does not change at run time. But when this gender information is available only dynamically (say from a file or database), we need an alternative mechanism to build the set clauses for the relation. Here we have relation `gender`, from section 1, which statically defines all the information.

```
// statically defined relation
relation gender(lref<string> p, lref<string> g)
{
  return eq(p,"Frank")  && eq(g,"male")
      || eq(p,"Sam")    && eq(g,"male")
      || eq(p,"Mary")   && eq(g,"female")
      || eq(p,"Denise") && eq(g,"female");
}
```

Let us assume this information regarding each person's gender has been read from a file or database into a `list<pair<string,string>>` called `genderList`. Let the first item in the pair be a name and the second be his gender. We can now define relation `gender_dyn`, based on `genderList` as follows:

```
list<pair<string,string> > genderList = ...;

// dynamically building a relation
Disjunctions gender_dyn(lref<string> p
                        , lref<string> g)
{
  Disjunctions result;
  list<pair<string,string> >::iterator i;
  for( i=genderList.begin();
                   i!=genderList.end(); ++i)
    result.push_back(
          eq(p,i->first) && eq(g,i->second) );
  return result;
}
```

Here we use the type `Disjunctions` to dynamically build the set of OR clauses for the relation. Type `Disjunctions` is itself a relation that supports dynamic addition of clauses. Thus we can trigger evaluation on it using operator(). The return type of `gender_dyn` has also been changed from relation to `Disjunctions`. This is optional, but useful as it implicitly conveys the dynamic nature of `gender_dyn` to consumers. Conceptually `Disjunctions` is simply collection of relations. When a `Disjunctions` instance is evaluated, the relations contained in it are treated as if there exists an `operator ||` between each pair of adjacent relations. Relations can be added at the front or at the back of a `Disjunctions` relation using

methods `push_back` and `push_front`. During evaluation, the contained relations are evaluated from front to back. Notice how an entire relational expression `eq(p,i->first) && eq(g,i->second)` is added to the `Disjunctions`.

Relations `Conjunctions` and `ExDisjunctions` are also provided for building relations dynamically. They correspond to operators `&&` and `^` respectively. The relation `gender` is redefined below using `Conjunctions` and `Disjunctions`.

```
Disjunctions gender_dyn(lref<string> p
                , lref<string> g) {
  Disjunctions result;
  Conjunctions conj1 = eq(p,"Frank");
  conj1.push_back( eq(g,"male") );

  Conjunctions conj2 = eq(p,"Sam");
  conj2.push_back( eq(g,"male") );

  Conjunctions conj3 = eq(p,"Mary");
  conj3.push_back( eq(g,"female") );

  Conjunctions conj4 = eq(p,"Denise");
  conj4.push_back( eq(g,"female") );

  result.push_back( conj1 );
  result.push_back( conj2 );
  result.push_back( conj3 );
  result.push_back( conj4 );

  return result;
}
```

In summary, `Conjunctions`, `Disjunctions` and `ExDisjunctions` together provide a facility for defining relations dynamically. This ability also naturally makes them a facility for runtime metaprogramming in the Logic paradigm.

2.7   Inline Logic Reference Expressions (ILE)

One may often encounter cases where the relations are simple operations to be performed on logic references in a relation simple expression application of standard operators like +, − etc. For instance, assuming the existence of a relation `multiply` for querying/asserting the product of two numbers, we can obtain the cube of a number as:

```
lref<int> n=2, sq, cu;
multiply(sq,n,n) && multiply(cu,sq,n)( );
cout << *cu ;
```

Using `multiply` the square is generated and from the square we obtain the cube. How would the above code look for computing the expression n*5-n/2+5? It is evident that writing out simple arithmetic expressions involving more than a couple operators gets verbose and unreadable quickly. Castor allows inline specification of expressions composed from standard operators and at least one logic reference. The cube example can now be rewritten more concisely:

```
lref<int> cu, n =2;
eq_f(cu, n*n*n)();
cout << *cu;
```

Relation `eq_f` unifies its first argument with the result of *evaluating* its second argument which is a function object. It is important to note that the value of `n*n*n` is not computed when it is passed to `eq_f`. Since the expression is composed using an lref it is turned into a function object and then passed to `eq_f`. The function object undergoes evaluation when the evaluation of `eq_f` kicks in. At the time of evaluation, `n` must be initialized, as accessing the value of an uninitialized lref results in an exception. Arbitrary expressions, such as n*5-n/2+5, composed of lrefs and common overloadable operators can be used to construct function objects easily:

```
eq_f(cu, n*5-n/2+5)();
```

Such inline specification of expressions involving logic references is called ILE, short for "Inline Logic reference Expression". ILEs come in handy in a variety of situations. Printing to console with relation `write_f` is another example where they can be used. The following example demonstrates the use of ILE for printing two strings separated by a comma:

```
lref<string> ls="Hello"; string s="world";
write_f(ls + string(",") + s)();
```

Relation `write_f` takes a function or function object as argument and prints the result of the evaluating its argument to stdout. We use an ILE in the above example to conveniently instantiate a function object and pass it to `write_f`. ILEs can be used as arguments to any relation provided by Castor suffixed with _f.

If T is the result type of an ILE, then all overloadable operators defined over T other than comma, dereference operator *, &, and -> can be used on `lref<T>` to produce an ILE. Currently ILEs lack support for mixing in objects of an arbitrary type T2 even if relevant operators are defined over T and T2. Thus,

```
write_f(ls + ",")(); //compiler error
```

fails even though operator + is defined for arguments of type `string` and `const char*`.

So far we have used ILEs to easily pass expressions to other relations. ILEs that produce a boolean value can be used to create simple relations inline. Consider the following relation which prints the result of the comparison of its arguments.

```
relation greaterLessEq(lref<int> n
                     , lref<int> cmpVal) {
  return write(n) && write(" is ") &&
   (  predicate(n<cmpVal) && write("lesser")
    || predicate(n>cmpVal) && write("greater")
    || write("equal")  );
}
```

Here, the ILE expressions `predicate(n<cmpVal)` and `predicate(n>cmpVal)` are used to conveniently define comparison relations directly inline instead of a defining global `less` or `greater` relations to do the same job. Similarly, `predicate(n%2==0)` can be used to create an inline relation to test for even numbers, or one could conceive more complex expressions like `predicate(n*2 >= cmpVal*n/2)`. It is important to note that relations defined using ILEs do not have the ability generate solutions. They will only assert if the said condition is true/false. Thus, `predicate(n<cmpVal)` will not generate values for `n` or `cmpVal`. All logic references involved in a relation are required to be initialized at the time when evaluation of the ILE occurs; otherwise an exception will be thrown. The exception will be thrown directly by any uninitialized logic reference when the ILE attempts to dereference it for evaluation. A more detailed discussion on ILEs is presented in section 4.

2.8 Sequences

Castor provides facilities for working with standard C++ sequence containers and iterators in a relational fashion. Tasks associated with sequences can be broadly categorized into producing sequences and iterating over the elements in it (i.e. consuming). Often, in traditional logic programming (as in functional programming or C++ template metaprogramming), operations requiring any modifications to a sequence are performed by creating a new sequence reflecting the necessary changes. Deletion of elements in a sequence is done by producing a new sequence without the unwanted elements. Thus deletion of elements is a combination of iteration over the original sequence and producing another sequence. Addition of new elements can also be performed in a similar fashion. Modifying values of elements can be considered an operation on the element and not on the sequence itself.

Although it is possible to write relations where deletion and insertion of elements is reflected directly in the original sequence, we will restrict ourselves to traditional logic style techniques. The following sections describe some of the common tasks surrounding sequences.

2.8.1 Generating Sequences

Sequences are created using the `sequence` relation. The following defines a relational expression that creates a list of three even numbers.

```
lref<list<int> > le;
relation evens = sequence(le)(2)(4)(6);
// see what it generates
if(evens())
  copy(le->begin(), le->end()
      , ostream_iterator<int>(cout, " "));
```

The uninitialized logic reference `le` is passed as the first argument to relation sequence. The elements used for

constructing the sequence are then consecutively passed individually to whatever is returned by the preceding function call. Since `le` is not initialized to a `list<int>` object, when the `sequence` relation is evaluated in the condition of the `if` statement, `le` will be initialized to a `list<int>` containing values 2, 4 and 6. Note that the `sequence` relation automatically figures out (from its first argument) if you are interested in creating a `list<int>` or `vector<string>` or some other sequence type. This ability is useful when writing generic code.

Abilities of the `sequence` relation go further. Sequences can be created not just out of simple values, but also from other sequences, logic references or an arbitrary mix of these. In the following code we create a list of strings using a mix

```
string s = "One";
lref<string> lrs = "Two";
vector<string> ls;
ls.push_back("Three"); ls.push_back("Four");
vector<string> lsTemp;
lsTemp.push_back("Five");
lsTemp.push_back("Six");
lref<vector<string> > lrls = lsTemp;
// create the sequence into ln
lref<vector<string> > ln;
relation numbers =
  sequence(ln)("Zero")(s)(lrs)(ls)(lrls);
// see what it generates
if(numbers())
  copy(ln->begin(), ln->end()
  , ostream_iterator<string>(cout," "));
```

A current limitation when creating a sequence from other sequences, all sequences involved should be of the same kind. That is, `lref<list<int> >` cannot be created directly from a `vector<int>` or `lref<vector<int> >` and vice versa. However this limitation is easily circumvented using `sequence`'s support iterators as follows:

```
vector<int> vi;
vector<int>::iterator b1, e1;
lref<list<int> > lrli1;
// generate using iterators to vi
relation r =
  sequence(lrli1)(vi.begin(), vi.end());


lref<vector<int> > lrvi;
lref<list<int> > lrli2;
lref<vector<int>::iterator> b2, e2;
// generate using logic reference iterators to
lrvs
relation r = begin(b2, lrvi) && end(e2, lrvi)
             && sequence(lrli)(b2,e2);
```

The relations `begin` and `end` will produce logical references that point respectively to the beginning and one past end of the sequence referenced by `lrvs`. The iterators are produced in a lazy fashion, i.e. only if and when they are actually evaluated in

the future. The logic references produced by `begin` and `end` are subsequently provided to `sequence` for constructing `lrls`. In general, care must be taken to avoid performing eager evaluations on logic references as they are typically initialized with appropriate values when backtracking and unification occurs.

2.8.2 Iterating over sequences

Relations `head`, `tail`, `next` and `prev` are provided for iterating through sequences. Relations `head` and `tail` provide relational style iteration that is also similar to how iteration is performed in functional languages. Relations `next` and `prev` provide support for iterating with iterators that and is closer to iteration techniques in traditional C++.

2.8.2.1   Iterating with head and tail

In this technique to iterate over all elements in a sequence container (like a list or vector), we split it into its head (i.e. the first element) and tail (i.e. collection of remaining elements). Splitting the tail again into its head and tail produces the second element in the original container. In this fashion we can continue to split the tail recursively till the tail is empty. The following example demonstrates the use of relations `head` and `tail` for printing all values in a list of integers.

```
relation printList(lref<list<int> > li) {
  lref<int> h;
  lref<list<int> > t;
  return head(li, h) && write(h) && write(",")
           && tail(li, t)
           && recurse(printList,t);
}
list<int> li;
// .. add numbers to list ..
printList(li)();//print elements
```

Relation `head_tail` is available for conveniently obtaining the head and tail directly in a single call. The return statement above can be rewritten using `head_tail` as follows:

```
return head_tail(li,h,t)
         && write(h) && write(",")
         && recurse(printList,t);
```

In cases where the tail is only used if a certain condition is true, then relations `head` and `tail` can be used more efficiently by computing the tail after evaluating the condition. Since computing a tail is an O(n) operation it makes sense to delay computing the tail after determining that it will be needed. A typical example of this is when searching for a vale in a list. Once the value of interest is found, the remainder of the list does not need to be processed. In cases like `printList`, where it is obvious that tail is always processed, `head_tail` may be preferred for sake of brevity.

Similar to `head` and `tail,` there also exists `head_n` and `tail_n` when the first or last *n* elements in the sequence are of interest. The following example generates the first and last 2 items from a vector containing 4 items.

```
int a[] = { 1,2,3,4 };
vector<int> v (a+0, a+4);

lref<vector<int> > h;
lref<vector<int> > t;

head_n(v, 2, h)();
// now h contains {1,2}
tail_n(v, 2, t)();
// now t contains {3,4}
```

The extra argument is used to specify the number of head/tail items we are interested in. Exception will be thrown if the size of the sequence is less than the specified head/tail size.

2.8.2.2   Iterating with next and prev:

Relation `next` represents a binary relation between a value and its successor. Since the successor of an iterator/pointer is another iterator/pointer that points to the next element in the sequence, we can use `next` to perform iteration. The following simple example prints the successor of 2:

```
lref<int> s;
next(2,s)();
cout << *n;
```

`next` can also be used to generate a predecessor:

```
lref<int> p;
next(p,2)();
cout << *p;
```

When both arguments are defined `next` will assert if the second argument is a successor of the first. An exception will be thrown if both arguments are undefined. Similar to `next`, relation `prev` is also available. The only difference between the two is that the order of arguments is reversed. `next` and `prev` can be used interchangeably depending on programmer's preference of which one is more readable in a given context.

The following example demonstrates use of `next` to print all items bounded by a pair of iterators.

```
relation printAll( lref<int*> beg_
                 , lref<int*> end_ ) {
  lref<int> val; // for storing **beg_
  lref<int*> n;
  return predicate(beg_==end_)
      ^ ( dereference(beg_, val)
         && write(val)
          && next(beg_,n)
          && recurse(&printAll,n,end_) );
}
```

```
int ai[]={1, 2, 3};
printAll(ai+0, ai+3)();
```

Relation `printAll` traverses through the items by recursively producing the successor to iterator `beg`. In each recursive step a check is performed to ensure that sequence is not empty by comparing `beg` to `end`. If there are elements in the sequence, we dereference `beg` (in a relational manner using `dereference`) to produce the underlying value and print it using `write`. After printing the first item, we proceed to recursing on the remainder of the sequence by producing successor of `beg` in `n`. The above example can also be rewritten in terms of iterators from the standard library, such as `vector<int>::iterator,` instead of simple pointer based iterators, by simply substituting occurrences of `int*` with `vector<int>::iterator` in the above example.

2.8.2.3   Iterating with item

Relations `next` and `prev` are useful in cases when explicit control over the iteration process is required. In cases when the intent is to simply produce all the values in a sequence one by one, relation `item` provides a simpler alternative. `item` takes 3 arguments; the first two are a pair of iterators (or a pair logic references to iterators) representing the sequence and the third argument is logic reference. An element from the sequence is produced in the third argument for each evaluation of `item`. The first two arguments must be defined.

```
int ai[]  = { 1, 2, 3, 4 };
vector<int> vi(ai+0, ai+4);
lref<int> val;
// 1 - iterating with regular iterators
relation r = item(vi.begin(), vi.end(), val);
while(r())
    cout << *val << ",";

// 2 - iterating with logic references to
//      iterators
lref<vector<int>::iterator> lBeg = vi.begin()
                    , lEnd = vi.end();
r = item(lBeg, lEnd, val);
while(r())
    cout << *val << ",";
```

In the above code, the third argument is left undefined in order to produce values from the sequence. `item` can also be used to assert if a particular value is present in the sequence simply by defining the third argument to the value of interest:

```
if( item(vi.begin(),vi.end(),4)() )
    cout << "found!";
```

2.8.3 Unification of Collections

Unification facilities provided by relation `eq` as discussed in section 2.3 above is not limited to scalar value types. Unification can be performed on collection types in a similar fashion.

```cpp
// produce a sequence
int a[] = { 1,2,3,4 };
vector<int> v (a+0, a+4);
lref<vector<int> > lrv;
eq(lrv,v)();
assert(*lrv==v);

// compare items
lref<list<int> > lrl= list<int>(a+0,a+4);
if(eq(lrl,v)())
  cout << "*lrl and v are equal" ;
```

Relation `eq` provides the basic unification support for collections (i.e sequences). In section 2.8.1 we discussed how `sequence` can be used to generate sequences. When the first argument is not initialized, `sequence` generates a sequence filled with the elements specified in the remaining arguments. However, if the first argument is initilized, it will perform comparison of the elements in the first argument with the items from the remaining arguments. This makes `sequence` a powerful unification facility for sequences. We have already seen examples of producing sequences in 2.7.1. The following example demonstrates its use for comparison.

```cpp
int a[] = {1,2,3};
lref<list<int> > lrl= list<int>(a+0,a+3);
lref<int> lri=1;
vector<int> v; v.push_back(2);
relation r =
    sequence(lrl)(lri)(v.begin(), v.end())
            (3);
assert(r()); // evaluation fails
```

Unlike `eq`, `sequence` allows creation/comparison of sequences using any mix of other sequences, iterators, lrefs of other sequences etc.

2.8.4 Summary

There are conceivably many ways of working with collections/sequences and Castor only provides support for a few useful ones. We can deal with collections directly as a whole or via itereators. Relations `item`, `next` and `prev`, `deref` are ligtweight facilites which deal with collections using iterators. Relations `eq`, `sequence`, `head`, `tail`, `item` allow working with whole collections directly. `sequence` is the most feature rich and flexibile, but is also a heayweight compared to the alternatives. Other relations such as `empty`, `size`, `insert`, `merge`, `eq_seq` etc. are also available for working with collections and iterators. may also be usedRefer to the reference manual for a complete list.

2.9   Cuts – Pruning alternatives

It is sometimes useful to discard paths that will not produce solutions or produce duplicate solutions when backtracking occurs. Such explicit pruning of paths is often done for efficiency reasons. There is no reason to waste time on pursuing alternatives paths that are known to fail or not produce anything of interest (e.g. duplicates of previously found solutions). Consider the following relation that searches through a binary search tree for a given value:

```cpp
// Binary search tree
struct BST {
  BST* l;     // left subtree
  BST* r;     // right subtree
  int value; // current node
   ...
};

relation b_search(lref<int> val
                 , const BST* tree) {
return
    predicate(val==tree->value)
 || predicate(val<tree->value)
      && recurse(b_search, val, tree->l)
 || predicate(val>tree->val)
      && recurse(b_search, val, tree->r)
   ;
}
```

Relation `b_search` consists of three clauses. The first clauses simply checks if the value of the current node matches the argument. The remaining two clauses will search recursively on the left and right sub trees depending on how `val` compares with the current node. If the first equality comparison succeeds, it is evident that the remaining clauses can be discarded. Once any of the comparison operations succeed we would like the backtracking mechanism to stay committed to the current clause and ignore other alternatives. The facility for eliminating alternative paths is typically called a *cut* in logic programming. Class `cut` and relation `cutexpr` provide support for cuts in Castor. We can rewrite the above relation using cuts as follows:

```cpp
relation b_search(lref<int> val
                 , const BST* tree) {
return cutexpr(
   predicate(val==tree->val) && cut()
|| predicate(val<tree->val)  && cut()
      && recurse(b_search, val, tree->l)
|| predicate(val>tree->val)
      && recurse(b_search, val, tree->r)
 );
}
```

Each occurrence of `cut()` marks the point at which we decide to commit to the current clause. The location of a `cut()` is called a *cut point*. A `cutexpr` only marks the boundaries within which pruning of alternatives takes place. In this case `cutexpr(…)` spans the all three clauses in the relation. Once the execution reaches one of the cut points, it will stay committed to the path beginning from the start of the `cutexpr` to the cut point. Thus all other alternatives available after "`cutexpr(`" begins will be discarded (or cut out) from consideration by the backtracking mechanism. Cuts do not affect the alternatives that existed prior to the `cutexpr` or the alternatives that exist after a cut point. The

cutexpr itself merely provides the scope or the extent within which the cut operation takes place.

A cut point without a surrounding `cutexpr`, or a `cutexpr` without any cut points are both meaningless. Such occurrences can sometimes occur accidentally when removing cuts from a relation that currently makes use of cuts. By design, such mismatched occurrences will produce compilation errors. Following usage of cuts wherein a `cutexpr(…)` appears in the caller and a `cut()` appears in the callee is also not allowed:

```
// Error: cannot dynamically nest cuts
relation outer(...) {
   return cutexpr( inner(..) || ... );
}

relation inner() {
   return  ... && cut() ...
}
```

Excessive usage of cuts in logic programming is generally discouraged since they tend to make logic programs less readable. Also, when not used with care, they can incorrectly prune out valid paths. Its usage should be preferred primarily in situations where it leads to a reasonable gain in performance. It is desirable for a relation that uses cuts, to also produce the same results when the cuts are removed. Cuts that do not alter the results of a relation are referred to as *green cuts*. Cuts that are not green are called *red cuts*. Many usages of cuts can be rephrased more elegantly using the relational ex-or operator.

2.10 Relational Ex-Or operator

We have seen operators && and || used to define relations. Castor also provides support for defining ex-or semantics between clauses in a relation using the operator ^. It is useful in expressing the idea that the second clause should be attempted only if first clause fails. Let us revisit the `greaterLessEq` example.

```
relation greaterLessEq(lref<int> n
                 , lref<int> cmpVal) {
 return predicate(n<cmpVal) && write("n<cmpVal")
     || predicate(n>cmpVal) && write("n>cmpVal")
     || write("n==cmpVal");
}
```

The following usage of `greaterLessEq` exposes a problem in the above implementation.

```
relation r = greaterLessEq(2,3);
while(r());
```

This produces the following output:
```
   n<cmpVal
   n==cmpVal
```

This is because the `while` loop forces the backtracking to purse all remaining paths even after the first clauses matches. The second clause fails due to the `(n<cmpVal)` guard at its

beginning and this leads to the evaluation of the third and final clause. Since the third clause does not have the guard `(n==cmpVal)`, it succeeds and thus we observe the output "`2 is equal`". An obvious solution is to put the guard in front of the third clause. Another solution is to introduce cut points at the end of the first and second clauses. However a simpler and preferred solution is to rewrite it in terms of the ^ operator as follows:

```
relation greaterLessEq(lref<int> n
                   , lref<int> cmpVal) {
return ( predicate(n<cmpVal)
                   && write("n<cmpVal") )
    ^ ( predicate(n>cmpVal)
                   && write("n>cmpVal") )
    ^ ( write("n==cmpVal") );
}
```

Here we are being explicit about the fact that three clauses are mutually exclusive candidates. Not only is this more explicit but also more readable and efficient. Once any one of the clauses succeeds `greaterLessEq` also succeeds. Backtracking will ignore any remaining unevaluated clauses. Thus all future attempts to seek more solutions from `greaterLessEq` will fail immediately. It is important to note the use of additional brackets around the clauses separated by ^. Since, unfortunately, the precedence of the ^ operator is higher than && and || operators, use of these brackets is necessary to preserve correct associativity.

Most, but not all, common usages of cuts can be elegantly replaced with use of ^ operator. Support for operator ^ to define relations is a feature unique to Castor.

2.11 Specifying Lref parameters types for relations

There are basically three ways of specifying lref parameters for relations:
- Basic: lref<T>
- By reference: lref<T>&
- By const reference: const lref<T>&

Basic: If the parameter type is specified to be `lref<T>`, arguments of type `T` and `lref<T>` are both acceptable. This versatility makes this mechanism a common choice for specifying parameters on relations. When the caller passes an argument of type `lref<T>`, the callee receives it such that the new lref points to the same underlying object of type `T`. This is similar to passing pointers in C++. However when the argument is an object `t` of type `T`, the callee receives a copy (on the heap) of `t` to which the callee's lref refers to.

```
relation foo(lref<int> i) {
    return True();
}

lref<int> li=2;
foo(1);  // OK! passes a copy of 1
```

```
foo(li);  // OK!
```

This basic parameter specification style is used most commonly in Castor.

By reference: For certain types where copy construction can be potentially very expensive such as `std::vector`, it is desirable to prevent implicit and (possibly) repeated copy construction. This can be done by disabling the ability to pass arguments of type `T` directly and requiring only arguments of type `lref<T>`. This can be done by specifying the parameter type as `lref<T>&`.

```
relation foo(lref<int>& i) {
    return True();
}

lref<int> li=2;
foo(1);  // Error!
foo(li); // OK!
```

All relations in Castor that expect container types (such as `empty` and `size`) as arguments use this form of parameter specification.

By const reference: This mode of parameter specification is similar to the basic style with one significant difference. A `const lref<T>&` parameter cannot be passed as argument to another relation that specifies its lref parameter by reference. E.g.:

```
relation bar(lref<int>& j) {
    return  ...
}

relation foo(const lref<int>& i) {
    return bar(i); // Error!
}
```

Yet another, but minor, difference compared to the basic style is that the value of the lref cannot be modified *directly* inside the relations:

```
relation foo(const lref<int>& i) {
    i=2; // Error!
    return bar(i);
}
```

This is rarely an issue since direct mixing imperative code inside a declaratively specified relation should never be done in practice. The reason being, the imperative assignment of 2 to `i` here is performed when the relation is invoked as opposed to when the relation is evaluated.

2.12 Debugging

Unfortunately, in an imperative world, debugging relational code is not as straight forward as debugging imperative code. Debuggers for C++ are, for good reason, designed to debug

imperative code conveniently. In principle, since we let the computer figure out how to solve the problem, we should be oblivious to how evaluation takes place. However in practice it is useful to be able to peek into the evaluation as it progresses. This is important for verifying correctness of user defined relations especially when the expected output does not match what is observed.

For debugging logic programs some creative thinking is typically required. Since relations return expressions that will be lazy evaluated, it is not very interesting to place a breakpoint directly inside the body of a relation's definition. In fact, some of the relation's `lref` arguments may not even be initialized with a value at the point where the relation is called, but are likely to be initialized by the time the relation undergoes evaluation. Also most of the actual execution occurs within the function objects returned by operators `||` and `&&` and relation `eq` and debugging using breakpoints inside these library artifacts requires understanding of Castor's implementation details[2].

A simple and primitive technique is to insert print statements at relevant points to observe the progress. Castor provides a relation `write` for printing to `stdout` in relational manner. In the following example we add debug statements to the `ancestor` relation to observe the generation of values P in the second clause as an attempt to find a solution progresses.

```
relation ancestor_dbg(lref<string> A
                , lref<string> D) {
  lref<string> P;
  return child(D,A)
     || child(D,P) && write(P) && write(",")
           && recurse(ancestor, A, P);
}
```

The following attempt to find all ancestors of Sam:
```
lref<string> X;
relation a = ancestor_dbg(X,"Sam");
while(a())
   cout << " :" << *X << " is Sam's
ancestor\n";
```

produces the following output on screen:
```
    :Mary is Sam's ancestor
    :Frank is Sam's ancestor
    Mary,Frank, :Gary is Sam's ancestor
    Gary,
```

It can be observed from this output that the ancestors Frank and Mary are discovered without resort to recursion since they are direct parents of Sam. Once both parents of Sam are found, the printed output reveals that backtracking proceeds to evaluate the recursive clause looking for Mary's ancestors.

---

[2] It may be interesting to note that the basic foundation to support logic programming is implemented in Castor in only about 400 lines of code.

3 Implementing relations imperatively

The ability to easily reach out from any paradigm to another, as and when required, is essential in multiparadigm programming. So far, we have seen examples of how relations can be built on top of other relations and also how relations can be consumed by imperative code. In this section we complete this cycle by understanding how relations can consume non relational facilities. This involves using imperative techniques to implement the relation.

Relations may be defined either declaratively or imperatively. The examples seen so far in this document are cases of defining them declaratively since the programmer is not involved in providing the operational semantics for the relation. That is, the algorithms and data structures used to perform computations are not specified by the author of the relation.

Producing imperative definitions for relations is generally more work compared producing declarative definitions. Also declarative definitions are easier to read and get right. However, there can be situations which motivate a programmer to define relations imperatively. Typical reasons include:

- Interaction will low level (e.g. I/O, memory, drivers etc.) or other imperative facilities for which relational abstractions either do not exists or the ones available are inadequate.
- Finer grain control over execution to improve performance when necessary.
- Improving the ability to step through a relation's execution with a debugger.

Castor provides a simple base class `OneSolutionRelation` that may be used when implementing relations imperatively. This class is useful *only* for implementing relations that evaluate successfully at most once.

For purpose of this discussion, relations are classified into test only relations and generative relations.

3.1 With relation predicate

A test-only relation is one that does not modify any of its arguments and consequently does not induce any side effects into the system that need to be reverted during backtracking. Such a relation is useful in testing if a given condition is true. Generally, due to their very nature, test-only relations succeed only once at most. Relation `predicate` is useful for defining such relations. The actual test condition can be packaged into a regular function that returns `bool`, and then invoked via relation `predicate`. Consider defining a relation that checks if a given file exists.

```
bool fileExists_pred(string fileName) {
    if(/* file found on disk */)
        return true;
    return false;
}
```

```
relation file_exists (lref<string> fileName_)
{
    return predicate(fileExists_pred,fileName_);
}
```

This is essentially a two step process. First, the predicate function `fileExists_pred` contains the imperative code to be used for the relation. Second, relation predicate is used to convert the predicate function into a relation inside the "wrapper" relation `file_exists`. Note that the relation takes an `lref` argument whereas the function takes a `string` argument. Although for both the parameter types can be `lref<string>`, it is natural for a relation's parameter type to be an lref and that of an ordinary function to be a regular type. Relation `predicate` assumes that the target predicate function does not take an lref argument. Consequently, if `filename_` is a logic reference, it automatically dereferences `fileName_`, when invoking the predicate function `fileExists_pred`. If `filename_` is not a logic reference, it will be passed directly.

Also note the customary use of '_' at end of the parameter's name 'fileName_'. This indicates that it is not a bi-directional argument. Thus its value should be defined by the time the evaluation of the relation takes place. If it is not defined at this time, an `InvalidDeref` exception is thrown when `predicate` attempts to dereference `fileName_`.

3.2 Using class Coroutine

Every relation is actually an instance of some function object type. So far we have defined relations as functions and expressions. Since this syntactic approach is a declarative way of specifying the operational details of the underling function object, its type name and definition are not visible to us. In this section we shall see how to implement these function objects directly. This has the advantage that it gives full control over the computational steps involved and also the relation is easier to step through using a debugger. The downside, however, is that it requires more code to implement and also needs some extra care.

We have seen that lazy evaluation is key aspect of how relations operate. Each time a value is generated, they return back to the caller immediately. On the next invocation they resume from where the left off and once again generate the value and return back to the caller. This is unlike typical functions which return back to the caller after generating all values. Functions that support this *suspend-and-resume* behavior are called coroutines.

Since C++ does not natively provide support for defining coroutines, Castor provides the class `Coroutine` and four macros (`co_begin`, `co_end`, `co_yield`, `co_return`). Although these were designed for implementing relations as classes, one may use them completely outside the context of logic programming.

In Castor, a coroutine is implemented as a function object. To define this class, we derive it from `Coroutine` and implement the function call operator `bool operator()(void)` as follows:

```cpp
// relation to check or generate values in
a specified inclusive range
template<typename T>
class Range_r : public Coroutine {
    lref<T> val, min_, max_;
public:
    Range_r(lref<T> val, lref<T> min_,
lref<T> max_)
        : min_(min_), max_(max_), val(val)
    { }

    bool operator() () {
      rel_begin();
      ...
      rel_end();
    }
};
```

Note the use of macro `rel_begin` to start and macro `rel_end` to end the body of `operator()`. No statements should precede or follow these two macros in the method body. These two macros merely set up a switch statement spanning the definition of `operator()`.

We can now implement `operator()` as a coroutine:

```cpp
    bool operator() () {
      rel_begin();
      if(val.defined())
         rel_return(    ( *min_<*val &&
*val<*max_ )
                     || ( *min_==*val )
                     || ( *max_==*val ) );
      for(val=min_; (*val<*max_) ||
(*val==*max_); ++val.get())
         rel_yield(true);
      val.reset();
      rel_end();
    }
```

Notice the absence of `return` statements. These have replaced with calls to macros `rel_return` and `rel_yield`. Both macros return the evaluated value of their argument back to the caller.

Macro `rel_yield` indicates a point where execution is temporarily suspended and a `true`/`false` value produced by evaluation of its argument is returned back to the caller. Next

invocation of `operator()` will *resume* execution inside the method starting directly at this *yield* point. Invoking this macro essentially causes the coroutine to remember the point where execution had reached in the previous invocation. On the other hand, invoking `rel_return` indicates that no future resumption of execution is required from inside the method body, and the `true`/`false` value produced by evaluation of its argument is returned to the caller. All future invocations of `operator()` on this instance of the class will return `false` immediately.

Conceptually, `rel_yield` indicates *suspension* of execution and `rel_return` indicates *completion* of execution. Since, in logic programming, returning `false` to the caller is an indication by relations that their task is complete; `rel_yield(false)` also indicates completion. If the argument to `rel_yield` evaluates to `false`, all future invocations on `operator()` will return `false` immediately.

One important thing to keep in mind when implementing these coroutines is to avoid defining variables inside `operator()`, since their state will not persist across invocations. This is demonstrated in example 8 below, which is a more natural but incorrect way of implementing example 7.

Examples:
```cpp
// 1) Simplest coroutine that never
succeeds
struct Simple1 : public custom_relation {
  bool operator()(void) {
    rel_begin();
    rel_end();
  }
};

Simple1 r;
cout << boolalpha << r() << "\n"; //
prints false
cout << boolalpha << r() << "\n"; //
prints false

// 2) Coroutine that succeeds once, using
rel_yield
struct Simple2 : public custom_relation {
  bool operator()(void) {
    rel_begin();
    rel_yield(true);
    rel_end();
  }
};

Simple2 r;
cout << boolalpha << r() << "\n"; //
prints true
cout << boolalpha << r() << "\n"; //
prints false
```

```cpp
// 3) Succeeds once, using rel_return
struct Simple3 : public custom_relation {
  bool operator()(void) {
    rel_begin();
    rel_return(true);
    rel_end();
  }
};

Simple3 r;
cout << boolalpha << r() << "\n"; // prints
true
cout << boolalpha << r() << "\n"; // prints
false


// 4) Succeeds twice, using rel_yield
  bool operator()(void) {
    rel_begin();
    rel_yield(true);
    rel_yield(true);
    rel_end();
  }


// 5) Succeeds only once
  bool operator()(void) {
    rel_begin();
    rel_return(true);
    rel_return(true); // will never be
executed
    rel_end();
  }


// 6) Succeeds only once
  bool operator()(void) {
    rel_begin();
    rel_yield(false);
    rel_yield(true);  // will never be
executed
    rel_end();
  }

// 7) Succeeds 'n' times
class SimpleN : public custom_relation {
  int n, i;
public:
  SimpleN(int n) : n(n), i(0)
  {}

  bool operator()(void) {
    rel_begin();
    while(i++ < n)
      rel_yield(true);
    rel_end();
  }
```

```cpp
};

// 8) Incorrect way of implementing
relation SimpleN from above
class SimpleN : public custom_relation {
  int n;
public:
  SimpleN(int n) : n(n)
  {}

  bool operator()(void) {
    rel_begin();
    for(int i=0; i<n; ++i) // 'i' should
not be defined here
      rel_yield(true);
    rel_end();
  }
};


// Example 9 : Relation to test/generate
size of a specified string
class StrSize : public custom_relation {
    lref<string::size_type> sz;
    lref<string> str_;
public:
    // str_ is an input only parameter,
sz is in/out
    StrSize(lref<string> str_,
lref<string::size_type> sz)
        : sz(sz), str_(str_)
    { }

    bool operator() (void) {
      rel_begin();
      if(sz.defined())
         rel_return( *sz == str_->size()
);
      sz = str_->size();
      rel_yield(true);
      sz.reset();  // revert external
side effects
      rel_end();
    }
};

cout << boolalpha << StrSize("blah",4)();
// prints true

lref<string::size_type> sz;
StrSize("blah",sz)();
cout << *sz; // prints 4
```

Also refer to:
predicate

3.3 Generative relations with single solution(delete this section)

Unlike test-only relations, a generative relation may modify its arguments. This is typically the case when an lref argument is not defined. But when all lref arguments are defined, it behaves like a test-only relation. When the value of one of its arguments is modified, these changes need to be reverted when backtracking decides to pursue an alternate path of evaluation. Reverting typically only involves resetting the lrefs to the undefined state. The generative relation is responsible for reverting any side effects it induced.

Class `OneSolutionRelation` provides support for implementing relations that generate at most one solution. Relation `FileContent`, in the following example, can be used to either obtain the contents of a file into a string, or to match the file contents against a string.

```
class FileContent
 : public OneSolutionRelation<FileContent> {
  lref<string> fileName, content;
  bool contentChanged;
public:
  FileContent(lref<string> fileName_
              , lref<string> content)
       : content(content)
       , fileName(fileName_)
       , contentChanged(false)
  { }

  bool apply() {
    string fileText = /* read file */;
    if(content.defined())
      return *content==fileText;// Test
    else {
      contentChanged = true;
      //Generate. i.e. modify content
      content = fileText;
      return true;
    }
  }

  void revert() {
    if(contentChanged) {
    // revert content to undefined state only if
it was modified
      content.reset();
      contentChanged=false;
    }
  }
};


lref<string> fileText;
relation r = FileContent("c:/blah.txt"
                         , fileText);
while(r()) // r() will only succeed once
  cout << *fileText;
```

This relation obviously cannot succeed more than once. The relation is defined as a class that inherits from `OneSolutionRelation<FileContent>`. The generic base class is parameterized on the derived class implementing the relation. Class `FileContent` defines member functions `apply` and `revert`. The `apply` method performs the job of generating a solution or validating a condition. The `revert` method reverts any side effects induced by `apply`. The base class only defines one member function `bool operator()(void)` which is inherited by the `FileContent`. When evaluation is triggered on the relation, this inherited `operator()` is invoked when then simply invokes either the `apply` or `revert` methods. When evaluation is triggered for this first time, `apply` is called and the return value obtained from it becomes the return value of `operator()`. Thus if `apply` succeeds, so does the first evaluation of the relation. When evaluation is triggered for the second time, `revert` is called. However `revert` is called only if `apply` had succeeded previously. It is considered that a relation that failed evaluation does not leave any side effects behind. `apply` checks to see if any side effects need to be reverted and accordingly performs the necessary reverts. `revert` will not be called more than once, regardless of the number of times evaluation is triggered on the relation. Note that `revert` does not return a `bool` value like `apply`.

Since `OneSolutionRelation` is only intended for relations that generate at most one solution, all attempts to evaluate the relation following the first attempt will implicitly return `false`.

3.4 Generative relations with multiple solutions (delete this section)

Under the covers, all relations defined (declaratively or imperatively) using Castor are coroutines with return type `bool`. Relations that generate multiple solutions (i.e. succeed more than once) are the most general form of relations. Test only relations and relations that generate only one solution are basically just special cases. This section describes the concept of a coroutine, how to implement them in C++ and finally how to implement a relation as a coroutine. Implementing relations as coroutines does not need any support from Castor.

3.4.1 Coroutines

All languages based on the imperative paradigm support the idea of a subroutine. The basic idea behind a subroutine can be summarized as follows. When a calling routine (i.e. parent routine) invokes a subroutine (i.e. the callee), the caller remembers the point at which control was transferred to the callee. The callee simply starts execution from the beginning and returns control back to the caller. The subroutine does not

remember the point at which it transfers control back to the parent. The parent routine now resumes execution from where it left off. Each time control is transferred from the caller to the callee, the callee essentially starts execution from its first statement. Conceptually, this mode of operation creates a parent/child relationship between a callee and a caller. Thus the name "subroutine".

Instead of a parent-child relationship, there exists a sibling relationship between the two routines if the callee is a coroutine. Both, caller and callee, remember the points at which transfer of control occurs. When control is transferred from the caller to a callee for the first time, the callee begins execution from the first statement. When the callee decides to yield control back to the caller, it too remembers this point at which the control is transferred back to the caller. Next time the caller decides to invoke the same callee again, the callee resumes execution at the point where it had previously left off and eventually, again, yields control back to the caller.

Since C++ does not support coroutines directly we need to simulate them in C++. Using hypothetical syntax, let us define a trivial coroutine that returns a number in the specified range, one at a time. Note that the coroutine produces only a single value on each evaluation.

```cpp
// coroutine
coroutine<int> numsInRange(int min, int max) {
   for(int i=min; i<=max; ++i)
      yield return i;
   return -1; // -1 indicates completion
}

// caller
coroutine<int> twentyNos = numsInRange(1,20);
for( int i=twentyNos(); i!=-1; i=twentyNos() )
   cout << i << " ";
```

Above, the coroutine returns -1 to indicate to the caller that all results have been produced. This is a simple (but not general) solution to the problem of how a coroutine would report completion to the caller. However, it suffices as far as our discussion in concerned. Once the `return -1` statement is reached, all future attempts at evaluating the coroutine will return the value -1 to the caller. In other words, the execution *gets stuck* at that point. This is unlike the `yield return` statement from where execution *resumes*. We can consider `return -1` a shortcut for:

```cpp
   while(true)
     yield return -1;
```

3.4.2 Implementing coroutines

An obvious solution to implementing the `numsInRange` coroutine in C++ is to turn it into a function object. This function object would have to remember the arguments with which it was invoked and the point at which to continue on the next evaluation attempt.

To implement the coroutine as a function object, let us start with a basic function object that saves away the arguments and implements the function call operator containing the original body of the coroutine as-is:

```cpp
class GetNumsInRange  {
  int min, max;
public:
  GetNumsInRange(int min, int max)
      : min(min), max(max)
  { }

  public int operator()(void) {
    for(int i=min; i<=max; ++i)
      yield return i;
    return -1;
  }
};
```

Now we shall make incremental modifications to the above function object till we have valid C++ code with identical semantics.

Step 1: Promote variables, which require their values to persist across evaluations, to data members:

```cpp
class GetNumsInRange  {
  int min, max, i;
public:
  GetNumsInRange(int min, int max)
      : min(min), max(max), i(min)
  { }

  public int operator()(void) {
    for( ; i<=max; ++i)
      yield return i;
    return -1;
  }
};
```

Step 2: Identify entry points[3]:
An exit point is any point at which the coroutine returns control back to the caller. Basically any `return` or `yield return` statement is any exit point. An entry point is a location in the body of the coroutine from where execution can directly begin when the coroutine resumes. Every `yield return` exit point introduces a corresponding entry point elsewhere in the coroutine:

-   The first instruction in the function body.
-   Any point from which evaluation resumes after a yield return statement.

```cpp
class GetNumsInRange  {
  int min, max, i;
public:
  GetNumsInRange(int min, int max)
      : min(min), max(max), i(0)
  { }
```

---

[3] Subroutines have a single entry point and multiple exit points whereas coroutines have multiple entry and exit points.

```
  public int operator()(void) {
 ❶
    for( i=min; ❷ i<=max; ++i)
      yield return i;
 ❸
    return -1;
  }
};
```

This is the best way to get the

```
class GetNumsInRange  {
  int min, max, i;
public:
  GetNumsInRange(int min, int max)
      : min(min), max(max), i(min)
  { }

  public int operator()(void) {
    for( ; i<=max; ++i)
      yield return i;
    return -1;
  }
};
```

3.4.3 Implementing relations as coroutines

4 Inline Logic Reference Expressions

An ILE is an expression composed of at least one logic reference variable and most of the common overloadable operators. Unlike typical expressions in C++, ILEs do not undergo evaluation to produce a result immediately at the point of definition. Instead they produce an expression tree that represents the semantics of the expression. This expression tree can be evaluated to produce a result, at a later point in time, by the applying the function call operator. The following example demonstrates these basic semantics:

```
int i=0;  // plain old variable
i+1+1;     // simple expression: produces value 2

lref<int> li=2;  // logic reference
(li+i)*2;         // ILE: produces an expression
tree
```

Following example uses an ILE where a function is required.

```
template<typename Func>
void printResult( Func f ) {
   cout << f();
}

printResult( li*li / 2 ); // pass ILE as
argument
```

In the above code, the ILE undergoes evaluation inside the printResult when f() is evaluated. We can also return functions from other functions:

```
void runtests();
boost::function<int()> halfOf( lref<int> li )
{
    return li/2; // return an ILE
}

boost::function<int()> f = halfOf(4);
cout << f(); // ILE is evaluated here
```

ILEs provide a convenient way to create expressions that require delayed evaluation. The convenience comes from not having to package them explicitly inside a named function. Such expressions are also traditionally called *lambda expressions*. Every ILE contains a copy of all variables and values required to compute its expression. This set of variables and values is referred to as the ILE's *closure*. Since copying logic references is like copying pointers or references (i.e. the copy refers to the same underlying object) any changes to the object referred to by the original lref will be observable to the ILE. When multiple lrefs refer to the same object, the object will be kept alive until the reference count goes down to zero. This makes it safe to evaluate ILEs even after the termination of the scope in which the ILE was created. This safety is essential in getting the most out of the delayed evaluation semantics provided by ILEs. The following examples illustrate the closure semantics.

```
int x=2;
lref<int> lx=3;
boost::function<int()> ile = lx+x;
cout << ile();  // prints 5
x=1; // this will not observable to the ile as
it contains a copy of x
cout << ile();  // prints 5
lx=4 ;    // updating the lref will be
observable to the ile
cout << ile();  // prints 8
```

Since an ILE is an expression and not a function, it does not accept any arguments at the time of evaluation. It is a free standing expression that can be evaluated at anytime and as many times as needed. As we saw above, each evaluation may potentially produce a different result. Changes in the result could occur due to external code altering the object referred to by one of the logic references, as in the code above. A change in the result could also be due to the ILE inducing side effects on its own closure as follows:

```
lref<int> lx=3;
boost::function<int()> ile = ++lx;
cout << ile(); // prints 4
cout << ile(); // prints 5
cout << ile(); // prints 6
```

ILEs such as `++lx` which induce side effects are said to be *impure*. Those that don't, such as `x+2`, are *pure*. [4]

4.1   Creating relations from ILEs:

ILEs were originally devised as a quick and easy way to create anonymous relations directly inline from simple expressions instead of having to create named relations. ILEs that return `bool` values can be turned into relations by with the help of relation `predicate`. Relation `predicate` is an adapter that allows functions or function objects (with up to six arguments) that return `bool` to be treated as relations. Relations created by passing an ILE to `predicate` are referred to as ILE-based relations.

The following example demonstrates the use of ILE to generate even numbers.

```
//Print all even numbers in the inclusive range
1...20
lref<int> x;
relation evens = range<int>(x,1,20)
                 && predicate(x%2==0);
while(evens())
    cout << *x << " ";
```

Above, the ILE-based relation `predicate(x%2==0)` tests if `x` is even. We start out with `x` not being initialized. Relation `range` generates values for `x`, and since `x` is part of the ILE's closure, any changes to `x` will be visible to the ILE. This allows the ILE-based relation to test the values one by one as they are generated. This is an example of the classic generate-and-test pattern commonly used in logic programming. In both the examples relation `range` is used to generate values for logic references and the ILE-based relations filter out values that fail the check.

The following example demonstrates the use of ILE to generate Pythagoras triplets.

```
// Print all Pythagoras triplets less than
30
lref<int> x,y,z; int max=30;
relation pythTriplets =
     range<int>(x,1,max)
  && range<int>(y,1,max)
  && range<int>(z,1,max)
  && (   predicate(x*x+y*y==z*z)
      || predicate(z*z+y*y==x*x) );

while(pythTriplets())
    cout << *x << "," << *y
         << ","<< *z << "\n";
```

Here the first three uses of `range`, generate values for `x`, `y` and `z` in the range 1 through 30 and the two ILE based relations are used to check if `x`, `y` and `z` form a Pythagoras triplet.

Another common usage of ILEs is with relation `eq_f`, which takes a function or function object as its second argument. It unifies the first argument with the value obtained by evaluating the second argument. Thus, if `x` and `y` are lrefs, then `eq_f(x, y*3)` unifies `x` with the value obtained by evaluating the ILE `y*3`.

Impure ILEs should not be used when creating ILE-based relations. If such relations produce side effects, they will not be reverted during backtracking and will most likely lead to surprising results. Consider `predicate(++x<5)` which increments value of `x` each time it is evaluated. Here, `++x<5` is an impure ILE as it modifies value of `x`. Backtracking relies on reversal of side effects. When such side effects are desired they should be packaged into a named relation that ensures reversal of the side effects. Relation `predicate(++x<5)` can be rewritten safely as `inc(x) && predicate(x<5)` or `next(x,y) && predicate(y<5)`. The former increments `x`, and the latter unifies `y` with the incremented value of `x` (without effecting value of `x`). The latter style is generally preferred.

The advantage of ILE-based relations is the brevity, the downside however is that they are not as full featured as named relations. ILE-based relations can only perform tests on values produced by other relations; they are unable to generate solutions themselves. Thus all logic references involved in the ILE must be initialized at the time of evaluation.

4.2   Limitations of ILEs:

There are a couple of limitations to ILEs that users should be aware of. The first one is due to language limitations, and the second one is by design.

1) Inferred return type: A seemingly trivial question that one may ask is "What is the type of the result produced on evaluation of an ILE?". The answer of course is "It depends on the expression represented by the ILE". A more precise answer should be "Same as if the expression were rewritten with logic references substituted with their effective types". That is, the expression `x*y` should yield an object of the same type and value regardless of whether `x` and `y` are lrefs or plain old variables. Unfortunately this is not always true as it is not possible to programmatically compute the return type of an arbitrary expression or function call in C++. Then how does one determine the return type of an ILE? It is inferred using the following intuitive rules:

-   All comparison operators (`<`, `>=`, `==`, `!=` etc.) and operators `&&`, `||` and `!` have return type `bool`.
-   Return type of prefix `++` and `--` is `T&`, if `T` is the argument type.

---

[4] This is same as the classification of pure and impure functions in computer science.

- All other unary and binary operators are assumed to have return type same as the type of their first argument.

An easy way of determining the type is to simply observe whether or not the ILE as a whole represents a comparison operation. If the ILE finally yields the result of a comparison, then its return type is `bool`, otherwise the return type is determined by the first argument to operator that is evaluated last in the ILE. Here are some examples:

```
lref<int> x=3;
double y=1;
2.0 * x;     // double
x * 2.0;     // int
y*x;         // double
x*y;         // int
2*x + 3.1;   // int
3.1 + 2*x;   // double
5 < x;       // bool
x == x;      // bool
x < 3*x+5;   // bool
3*x+5 < x;   // bool
++x;         // int&
```

2) Supported operators: All overloadable operators with exception of the following are supported for the creation of ILEs.
- AddressOf operator `&`
- Dereference operator `*`
- Member access operator `->`
- Indexing operator `[]`
- Comma operator `,`
- All forms of assignment (`=`, `+=`, `*=` etc.)

Since assignment is used to assign a new value to an lref immediately, it obviously cannot be used to produce ILEs. Similarly, dereference operator `*` and operator `->` and are used to access the underlying value and it rules them out from being used to compose ILEs. AddressOf and comma operators take on default language defined semantics when used on lrefs. Finally, as there is no reasonably intuitive way to know the result type of an index operator, it too remains unsupported.

4.3   Summary:

The following is a gist of the major points covered in this section on ILEs:
- ILEs are expressions which can be treated as function objects.
- Every ILE carries its closure with it, thus it can be safely evaluated even after the termination of their lexical scope.
- ILEs are either pure (i.e. free of side effects) or impure.
- Only pure ILEs should be used to create ILE-based relations.
- ILE-based relations cannot generate solutions; they can only perform testing of values generated by other relations.

- Return types of ILEs is determined by applying the two rules noted above.
- Most overloadable operators are supported on lrefs for creating ILEs.

5   Limitations of Logic Paradigm

5.1   Bi-directionality of lref arguments.

Typically lref parameters for relations are bi-directional. A relation generates a value for an lref parameter if its value is left undefined by the consumer of the relation. However it is not always feasible or even sensible for a relation to generate values for some of its lref parameters. Values for these lrefs must be defined by the time the relation is evaluated. Such lref parameters are called *input-only* parameters.

Consider implementing relation `not_eq` which succeeds if the two arguments are not equal.

```
relation not_eq( lref<int> lhs_
                   , lref<int> rhs_)  {
    return …
}
```

In this case, it is straightforward to compare the two arguments and check if they are not equal. But if one of the arguments is not initialized, on what basis could we generate a sensible value(s) for such that it is not equal to the other? There are multiple possibilities but no single approach can be easily deemed to be any better than another. For instance, we could start at 0 and go on till we hit `numeric_limits<int>::max()`. Or we could start at `numeric_limits<int>::min()` continuing until `numeric_limits<int>::max()`. Perhaps even backwards starting from `numeric_limits<int>::max()`.

Attempting a search in such a wide spanning range is rarely a very productive or useful way to solve problems. The futility of such an approach will seem even more obvious if we consider defining `not_eq` for `string`. Thus it only makes sense for `lhs_` and `rhs_` to be input-only parameters. By convention, names of input-only parameters contain a '`_`' suffix, and in-out parameters do not.

Another such example is the relation `size` provided by Castor.

```
relation size(lref<Cont>& cont_
        , lref<typename Cont::size_type> sz) {
  return ...
}
```

The first argument is any standard container and the second argument represents the size of a container. It is straightforward to test if `cont_` has size `sz` and also easy to generate a value for `sz` given an arbitrary container object. But what `sz` is initialized and `cont_` is not? It makes no sense to come up with random containers having size `sz` with arbitrary elements. Thus we specify `cont_` as an input-only parameter.

In each of these cases, it is left up to the consumer of these relations to decide how the values for an input-only parameter is to be generated.

5.2   I/O is not reversible.

Backtracking by its very nature depends upon the ability to revert side effects when pursuing alternate paths of evaluation. I/O, on the other hand, is rarely something that can be reverted. Once some data has been written to a socket or a document has been printed, there is no going back. Thus I/O and backtracking are at odds with each other. There are two ways to make these co-exist. First option is demand support from all kinds of I/O to be reversible. This is clearly not practical. The second option is to turn a blind eye to any I/O that occurs during the evaluation of relations. This is the only practical way in which is the two can be made to co-exist. However, it is up to the consumer of any relation that performs I/O to ensure that such side effects do not impact or interfere with the choices made during backtracking. They should be used in a way that is "backtracking-safe".

Relations `write` and `read` are examples of I/O relations in Castor. Consider the following misuse of `read` to check if the word read from stdin is either "Hello" or "World".

```
relation r = read("Hello") ^ read("World");
```

Although the declarative reading "read either Hello or World" is sound, it assumes that if the first `read` fails then its side effects will be reverted so that the second `read` can re-read the data. The correct way to perform this operation is to call `read` only once as follows:

```
lref<string> w;
relation r = read(w) &&
       (eq(w,"Hello") ^ eq(w,"World"));
```

Here we read the value, whatever it may be, into `w`, and then check if `w` is either "Hello" or "World".

6   Logic Programming Examples

In this section we solve a few problems using LP to get a better feel for logic programming techniques.

6.1   Directed acyclic graphs.

Figure 1

Here we consider the problem of finding paths in the above directed acyclic graph consisting of 5 edges. The graph can be easily described in terms of its edges. In the following edge relation, each line represents a directed edge of the graph. The parameters n1 and n2 represent the start point and the end point of an edge respectively.

```
// Representation of the directed edges in the
DAG shown in figure $$
relation edge(lref<int> n1, lref<int> n2) {
  return eq(n1,1) && eq(n2,2)
      || eq(n1,2) && eq(n2,3)
      || eq(n1,3) && eq(n2,4)
      || eq(n1,5) && eq(n2,4)
      || eq(n1,5) && eq(n2,2)
  ;
}
```

Given the definition of an edge we say that a path exists between two nodes Start and End if:
  • A directed edge exists between Start and End. **OR**
  • A directed edge exists between Start and some node (nodeX) **AND** a path exists from nodeX to End.
This relation can thus be defines as follows:

```
// Rule defining the meaning of a path
relation path(lref<int> start, lref<int> end)
{
  lref<int> nodeX;
  return edge(start, end)
      || edge(start, nodeX)
          && recurse(path, nodeX, end);
}
```

Given the relations `edge` and `path` above, we are in a position to ask many different questions concerning paths in the directed acyclic graph. A simple question is to see if a path exists between any two nodes:

```
relation p1_4 = path(1, 4);
if(p1_4())
  cout << "Path exists from 1 to 4";
```

A typical problem in graphs is to find all nodes reachable from a particular node. The following code prints all nodes reachable from node 5.

```
lref<int> node;
relation p5 = path(5, node);
while(p5())
 cout<< *node <<" is reachable from 5\n";
```

By leaving variable `node` uninitialized above, we are letting the computer generate values for it. It is interesting to note that 4 is listed twice when the above code is executed. That is due to the fact that node 4 can be reached via two independent paths from node 5, and during the process of inferencing both paths are discovered. Similarly, leaving the first argument uninitialized and specifying 5 as the second argument, gives a list of all nodes from which node 5 can be reached.

It is also possible to print all pairs of nodes between which a path exists by simply leaving both arguments to path uninitialized:

```
lref<int> n1, n2;
relation p = path(n1, n2);
cout<< "Path found between these nodes\n"
while(p())
  cout << "(" << *n1 <<","<< *n2 << ") ";
```

The above code will print the following pairs in order: (1,2) (2,3) (3,4) (5,4) (5,2) (1,3) (1,4) (2,4) (5,3) (5,4). Once again the pair (5,4) is listed twice as two paths were found from 5 to 4.

6.2 Finite Automata (FA)

Finite automata can be represented in the form of a set of transitions and a set of final states. Class Nfa below encodes the nondeterministic finite automata (NFA) from figure 2 which is equivalent to the regular expression ((ab)*ba)|b. It has two static member relations `transition` and `final`. Relation `transition` describes each transition in terms of the two states and the input character that make up the transition.

RegEx: **(**(ab)*ba**)**|b
Start State: 0
Final States: 3, 4
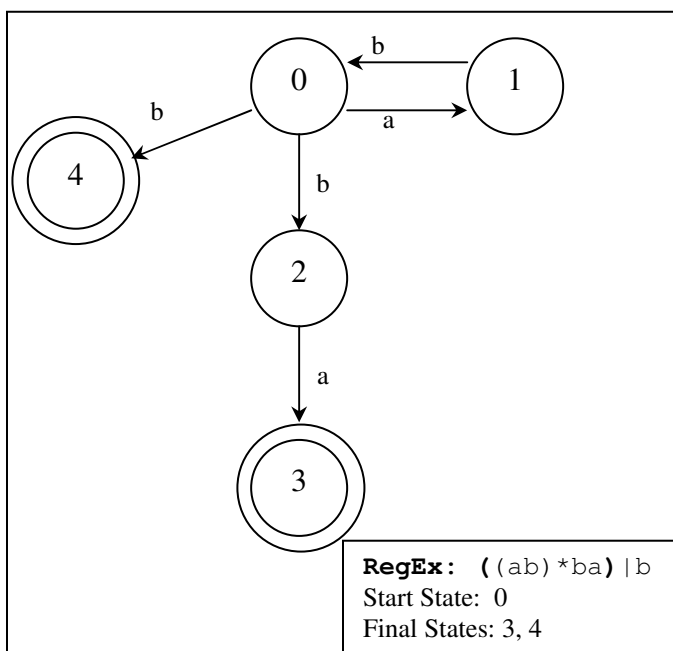
Relation `final` specifies the final states in the NFA.

A global relation `run` defines the rule for a successful run of any FA over an input string. The specific Finite Automata over which it operates is specified as a template parameter. The run defines that a run is successful if:
- Input string is empty **AND** a final state has been reached, **OR**
- There exists a transition from current state (`st1`) to some other state (`st2`) over the first character in the input string **AND** there is a successful run over the rest of the string starting from state `st2`.

Figure 2

Class NFA's static member relation `transition` describes the state transitions and member relation `final` describes the final states. Global relation `run` is takes NFA as a template parameter. It executes the NFA over an input string by matching one of the possible transitions to the next character.

```
class Nfa { // => ((ab)*ba)|b
  static const char a = 'a', b = 'b';
public:
  // tranistions NFA
  static relation transition( lref<int> st1
                            , lref<char> ch
                            , lref<int> st2) {
    return
       eq(st1,0) && eq(ch,a) && eq(st2,1)
    || eq(st1,0) && eq(ch,b) && eq(st2,2)
    || eq(st1,0) && eq(ch,b) && eq(st2,4)
    || eq(st1,1) && eq(ch,b) && eq(st2,0)
    || eq(st1,2) && eq(ch,a) && eq(st2,3)
    ;
  }

  // all final states of the NFA
  static relation final(lref<int> state) {
    return eq(state,3) || eq(state,4);
  }
};

// rule determines successful run of a FA
template<typename FA>
relation run(lref<string> input
            , lref<int> currSt=0) {
  lref<char>   firstCh;
  lref<string> rest;
  lref<int>    nextSt;

 return
   eq(input,"") && FA::final(startSt)
|| head(input,firstCh) && tail(input,rest)
     && FA::transition(currSt,firstCh,nextSt)
     && recurse(run<FA>, rest, nextSt)
   ;
}
```

Given the above relations, any input string can be tested against the finite automata as follows:

```
if( run<Nfa>("aabba")() )
```

```cpp
        cout << "Matched";
else
        cout << "No match";
```

6.3 Query Expressions

7 References

[CastorDesign] Roshan Naik, **Blending the Logic Paradigm into C++:** Discusses the design and implementation of the core of Castor that enables the LP in C++. http://www.mpprogramming.com/resources/CastorDesign.pdf

[CastorReference] Roshan Naik, **Castor reference manual:** http://www.mpprogramming.com/resources/CastorReference.pdf

[Leda] Timothy Budd, **Multiparadigm programming in Leda:** Addison-Wesley, 1995 : This book is a rich source of multiparadigm programming techniques. The presentation is the some seminal work related to multiparadigm language Leda and